

## Building an Intervention System Student Intervention

### 1. Classification vs Regression

*Your goal is to identify students who might need early intervention - which type of supervised machine learning problem is this, classification or regression? Why?*

This is a supervised classification prediction problem. Classification deals with a discrete set of output labels, such as 'pass' or 'fail'. The objective of a supervised classification problem is to learn about the relationship between inputs and discrete outputs (e.g. pass or fail) and then predict these discrete outputs on new data. Regression is not useful for this exercise as it is more suited for problems with continuous values, such as 5.9 or 9.4.

### 2. Exploring the Data

*Can you find out the following facts about the dataset?*

- Total number of students: 395
- Number of students who passed: 265
- Number of students who failed: 130
- Graduation rate of the class (%): 67.09
- Number of features: 30

*Use the code block provided in the template to compute these values.*

### 3. Preparing the Data

*Execute the following steps to prepare the data for modeling, training and testing:*

- Identify feature and target columns
- Preprocess feature columns
- Split data into training and test sets

*Starter code snippets for these steps have been provided in the template.*

## Building an Intervention System Student Intervention

### 4. Training and Evaluating Models

*Choose 3 supervised learning models that are available in scikit-learn, and appropriate for this problem.*

#### A. Decision Tree Classifier

- *What are the general applications of this model? What are its strengths and weaknesses?*

The decision tree classifier, as the name implies, is a classification algorithm. This classifier is relatively straightforward and easy to understand (e.g. it creates a downward tree structure). One of its most well known strengths is it can be visually graphed.

One of the decision tree's greatest weaknesses is that it is prone to overfitting out-of-the-box. Tweaking is required to stop the growth of the tree at the appropriate time. I've also learned that decision trees create biased trees if some classes dominate.

- *Given what you know about the data so far, why did you choose this model to apply?*

I chose this model as it's a simple model to understand and it handles both numerical and categorical data (although the categories for this data set have been transformed to 1's and 0's).

Looking at the data, I would think that the 'gini' criterion will identify the best feature to split the data early-on to create the shortest tree possible. I believe features, such as previous course failures (failures), will be split earlier within the decision tree process.

## Building an Intervention System Student Intervention

- Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.
- Produce a [table](#) showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.

DecisionTreeClassifier	Training set size		
	100	200	300
Training time (secs)	0.002	0.002	0.003
Prediction time (secs)	0.000	0.000	0.000
F1 score for training set	1.0	1.0	1.0
F1 score for test set	0.6611570247	0.7591240875	0.7441860465

### B. K Neighbors Classifier

- What are the general applications of this model? What are the general applications of this model? What are its strengths and weaknesses?

K Neighbors Classifier makes its prediction based on other points of data that are closest to the target. The “K” is the number of neighbouring labels to include when predicting the target label. The “K” is setup by the user. This algorithm is known as a lazy learner. In other words, it does all of it’s calculation not in training, but in querying/ predicting the target label.

- Given what you know about the data so far, why did you choose this model to apply?

I chose to test out this model as I feel that K-NN will do a good job in correctly labeling target students based on previously scored students with “neighbouring” data. My main concern is the processing power (memory) required during query time.

## Building an Intervention System Student Intervention

- *Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.*
- *Produce a [table](#) showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.*

KNeighborsClassifier	Training set size		
	100	200	300
Training time (secs)	0.000	0.001	0.001
Prediction time (secs)	0.002	0.005	0.009
F1 score for training set	0.8321167883	0.8501742160	0.8452655889
F1 score for test set	0.791176470	0.7714285714	0.8085106382

### C. Support Vector Classifier

- *What are the general applications of this model? What are its strengths and weaknesses?*

A Support Vector Classifier belongs to the Support Vector Machine family of algorithms. These sets of algorithms can be used for linear and non-linear classification problems. SVMs work best when there is a clear margin of separation between classes. A kernel function can be used to solve non-linear problems by transforming them into linear problems. This is said to be using higher-order parameters.

Support vector algorithms have been successful in many real-world problems such as text (and hypertext) categorization, image classification, bioinformatics (Protein classification, Cancer classification), and hand-written character recognition.

A support vector algorithms are known to work well when there is a clear margin of separation between classes. Another known strength is that the algorithm still works well when there is a high number of dimensions and the sample size is small. It also pays less attention to outliers as it only uses a subset of the training data.

Support vector algorithm may struggle when very large datasets are used or when there is a lot of pollution in the data (noise in the data). Support vector

## Building an Intervention System Student Intervention

algorithms are also known to be prone to overfitting out-of-the-box. Tweaking parameters are required to circumvent this challenge.

- *Given what you know about the data so far, why did you choose this model to apply?*

The SVC algorithm looks to be a good fit for the data as it's data is rather small in size. If any outliers exist within the data, SVC is designed to pay less attention to those outliers.

- *Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.*
- *Produce a [table](#) showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.*

SVC	Training set size		
	100	200	300
Training time (secs)	0.001	0.004	0.012
Prediction time (secs)	0.001	0.003	0.010
F1 score for training set	0.8951048951	0.878688524	0.8596112311
F1 score for test set	0.8129032258	0.8205128205	0.8235294117

## Building an Intervention System Student Intervention

### D. GaussianNB

- *What are the general applications of this model? What are its strengths and weaknesses?*

GaussianNB is part of the Naïve Bayes family of algorithms. Naïve Bayes algorithms have been used to efficiently identify spam emails and fraudulent activity on Visa cards. It is known to be very effective in text classification problems.. This algorithm is known to work well in large datasets with a lot of noise. It is fast to train and classify; this is because the algorithm assumes that features are independent of each other. It is not sensitive to irrelevant data.

Naïve Bayes algorithms are known to be weaker with smaller datasets. The algorithm is also known to assume independence of features; this is an asset when dealing with large and noisy data but it becomes a double-edged sword when dealing with a smaller and more predictable datasets.

- *Given what you know about the data so far, why did you choose this model to apply?*

Honestly, I chose a Naïve Bayes algorithm because I keep hearing it's the gold standard in machine learning. I feel like this algorithm had to be included in the top four of my choices for this exercise (although this assignment only asked for three algorithms).

But if I can try to answer this question, I would say that given the data, a Naïve Bayes algorithm would not be impacted by irrelevant features (if irrelevant features exist in the data). I also chose this algorithm in the case where the dataset would contain a lot of noise in the data; a Naïve Bayes algorithm still works well with noisy data.

## Building an Intervention System Student Intervention

- *Fit this model to the training data, try to predict labels (for both training and test sets), and measure the F1 score. Repeat this process with different training set sizes (100, 200, 300), keeping test set constant.*
- *Produce a [table](#) showing training time, prediction time, F1 score on training set and F1 score on test set, for each training set size.*

GaussianNB	Training set size		
	100	200	300
Training time (secs)	0.001	0.001	0.001
Prediction time (secs)	0.000	0.000	0.000
F1 score for training set	0.4337349397	0.8243727598	0.8243727598
F1 score for test set	0.4222222222	0.7883211678	0.7883211678

### 5. Choosing the Best Model

*Based on the experiments you performed earlier, in 1-2 paragraphs explain to the board of supervisors what single model you chose as the best model. Which model is generally the most appropriate based on the available data, limited resources, cost, and performance?*

In my analysis, I applied four algorithms to the dataset; Decision Tree Classifier, K Neighbors Classifier, Support Vector Classifier, and Gaussian Naïve Bayes. The model producing the best predictive score, F-1 score in technical terms, is the K neighbors classifier.

The K neighbors classifier algorithm produced the best predictive score based on the available data, limited resources, cost and performance. The K neighbors classifier is a stronger algorithm when the dataset was small, in this case, 395 records. This algorithm will perform well under limited resources because it is a lazy learner. In other words, a lazy learner algorithm does not ask for much resources when learning about the data, and it only requires a small amount of resources when querying a small dataset. I conclusion, The K neighbour classifier performs well with smaller data sizes, without asking too much on memory and other resources, resulting in an effective and low cost solution.

## Building an Intervention System Student Intervention

*In 1-2 paragraphs explain to the board of supervisors in layman's terms how the final model chosen is supposed to work (for example if you chose a decision tree or support vector machine, how does it make a prediction).*

The best model for predicting the pass or fail label for student data is the K Neighbors Classifier. This solution, based on the criteria of available data, limited resources, cost and performance, performed better than the other 3 algorithms chosen for this analysis: Decision Tree Classifier, Support Vector Classifier, and Gaussian Naïve Bayes.

So how does the K neighbors Classifier algorithm works you may ask? Similar to selling a house in a neighbourhood, with all else equal, the house will have the same value as the other houses around it. The “K” is the number of houses around the target you want to include in the equation. In my analysis, the best predictive score included the nearest 5 data points for each of the features to make the most accurate prediction for a student's pass or fail result.

*Fine-tune the model. Use gridsearch with at least one important parameter tuned and with at least 3 settings. Use the entire training set for this.*

```
parameters = {'n_neighbors':(4,5,6),'algorithm':('ball_tree','kd_tree','brute'),'p':(1,2)}
```

*What is the model's final F1 score?*

KNeighborsClassifier	Training set size: 300
	Testing set size: 95
Training time (secs)	0.226
Prediction time (secs)	0.004
F1 score for training set	0.8465116279
F1 score for test set	0.8057553956